

PyTorch fundamentals

	typename = torch.tensor()			
type	Scalar	vector	matrix	tensor
ex)	7	[7, 7]	[7, 8], [9, 10]	[[[1, 2, 3], [3, 6, 9], [2, 4, 5]]]
.ndim	0	1	2	3
.shape	[]	[2]	[2, 2]	[1, 3, 3]
.item()				
visualisation	7	$\begin{bmatrix} 7 \\ 4 \end{bmatrix}$	$\begin{bmatrix} 7 & 10 \\ 4 & 9 \end{bmatrix}$	$\begin{bmatrix} 7 & 4 & 6 & 1 \\ 1 & 9 & 2 & 3 \\ 5 & 6 & 8 & 8 \end{bmatrix}$
	torch.size()			

random_tensor = torch.rand(size=(r, c)) ... rand produces floats uniformly distributed across 0-1
 .zeros
 .ones

range_tensor = torch.arange(start, end, step)

tensor_like_zeros = torch.zeros_like(input = range_tensor) // creates tensor like range_tensor (same shape)

tensor datatypes: t = torch.tensor(, dtype=None, device=None, requires_grad=False)
 float32, float16, double, complex64, complex128
 "cpu", "cuda" -> where is tensor on?

conversion between datatypes: float16_tensor = float32_tensor.type(torch.float16)

getting info from tensors: .dtype .shape .device

tensor manipulation: tensor +/- * / 10 // operates 10 on each index
 torch.add(tensor, 10)

tensor * tensor // element wise multiplication

matrix multiplication torch.matmul(tensor, tensor)

transpose tensor.T

tensor aggregation: .min(), .max(), .mean(), .sum() // returns values
 .argmin(), .argmax() // returns index of values of tensor[0]

to fix dimension issues

a = y.reshape(x, y)

a = y.view(x, y) // returns view of original tensor in different shape but stores same data as original tensor

torch.stack([inputs], dim=0) // combine multiple tensors on top of each other or side by side

.squeeze() // removes all single dimensions from a target tensor

.unsqueeze(dim=x) // adds single dimension at dimension x

x = permute([new dimensions]) (2, 0, 1) // returns view of original input with dimensions permuted (reordered)

indexing: `x = torch.arange(1,10).reshape(1,3,3) ... x[0][0] = tensor([1,2,3])`
`...` will select all of a target dimension `... x[0][1][0] = tensor([1,4,7])`

PyTorch tensors & numpy: `torch.from_numpy(ndarray)` // numpy data to PyTorch tensor

`torch.Tensor.numpy()` // PyTorch tensor to numpy

(c) `import torch ... import numpy as np ... array = np.arange(1.0,8.0)` ... tensor = `torch.from_numpy(array).type(torch.FloatTensor)` convert to tensor

`print(torch.all_tensors)`

reproducibility: `RANDOM_SEED=42 ... torch.manual_seed(RANDOM_SEED)`

check for GPU access with PyTorch:

`import torch ... torch.cuda.is_available()`

Setup device agnostic code:

`device = "cuda" if ... else "cpu"`

Count number of devices:

`torch.cuda.device_count()`

Putting tensors and models on GPUs:

`tensor_on_gpu = tensor.to(device)` when instantiated, defaults to cpu

Moving back to CPU:

`tensor_back_on_cpu = tensor_on_gpu().numpy()` numpy needs cpu