## pytorch fundamentals

typename = torch.tensor( )

| type | scalar | vector | matrix | tensor |
|---|---|---|---|---|
| ex) | 7 | [7,7] | [7,8],[9,10] | [[[1,2,3],[3,6,9],[2,4,5]]] |
| .ndim | 0 | 1 | 2 | 3 |
| .shape | [ ] | [2] | [2,2] | [1,3,3] |
| .item() | | | | |
| visualization | 7 | $\begin{bmatrix}7\\4\end{bmatrix}$ [7 4] | $\begin{bmatrix}7 & 10\\4 & 3\\5 & \cdot\end{bmatrix}$ | $\begin{bmatrix}[7\ 4]\ [0\ 1]\\[1\ 9]\ [2\ 3]\\[5\ 6]\ [8\ 8]\end{bmatrix}$ |

torch.Size( )

random_tensor = torch.rand(size=(r,c)) ... rand produces floats uniformly distributed across 0-1
.zeros
.ones

range_tensor = torch.arange(start,end,step)

tensor_like_w_zeros = torch.zeros_like(input = range_tensor) //creates tensor like range_tensor (same shape)

tensor datatypes: t = torch.tensor( , dtype=None, device=None, requires_grad=False)
(float32, float16; tracks gradients yes or no?; "cpu", "cuda" what device is tensor on?)

conversion between datatypes: float_16_tensor = float_32_tensor.type(torch.float16)

getting info from tensors: .dtype .shape .device

tensor manipulation: tensor +/-/*/ 10 //operates 10 on each index

torch.add(tensor, 10)

tensor * tensor //element wise multiplication

matrix multiplication: torch.matmul(tensor, tensor) (.mm)

transpose: tensor.T

tensor aggregation: .min() .max() .mean() .sum() //returns values (requires tensor of float 32)

.argmin() .argmax() //returns index of values of tensor[0]

to fix dimension issues:

a = b.reshape(x,y) (will work if a has x*y elements)

a = b.view(x,y) //returns view of original tensor in different shape but shares same data as original tensor

torch.stack([inputs], dim= ) //combines multiple tensors on top of each other or side by side

.squeeze() //removes all single dimensions from a target tensor

.unsqueeze(dim=x) //adds single dimension at dimension x

x.permute(new dimensions of) (2,0,1) (1st dimension first, 0th dimension 2nd, 2nd dimension 3rd) //returns view of original input with dimensions permuted (rearranged)

indexing: ex) x = torch.arange(1,10).reshape(1,3,3) ... x[0][0] = tensor([1,2,3])

: will select all of a target dimension ... x[0][:][0] = tensor([1,4,7])

pytorch tensors & numpy: torch.from_numpy(ndarray) // numpy array → pytorch tensor

torch.Tensor.numpy() // pytorch tensor → numpy

ex) import torch ... import numpy as np ... array = np.arange(1.0, 8.0) (default is float64) ... tensor = torch.from_numpy(array).type(torch.float32) (want it as tensor)

reproducibility: .min(tensorA == tensorB) ——— RANDOM_SEED = 42 ... torch.manual_seed(RANDOM_SEED)

check for gpu access with pytorch: import torch ... torch.cuda.is_available()

setup device agnostic code: device = "cuda" if ↓ else "cpu"

count number of devices: torch.cuda.device_count()

putting tensors and models on gpu's: tensor_on_gpu = tensor.to(device) (when instantiated, defaults to cpu)

↳ moving back to cpu: tensor_back_on_cpu = tensor_on_gpu().numpy() (numpy needs cpu)