pytorch paper replicating

0: getting setup ## need torch 1.12+ and torchvision 0.13+

import torch ... import torchvision ... from torch import nn ... from torchvision import transforms ... from torchinfo import summary

from going_modular.going_modular import data_setup, engine

from helper_functions import download_data, set_seeds, plot_loss_curves

try and exception blocks if uninstalled

device = "cuda" if torch.cuda.is_available() else "cpu"

1: get data

image_path = download_data(source = " github raw link ", destination = "pizza_steak_sushi")

train_dir = image_path / "train" ... test_dir = image_path / "test"

2: create datasets and dataloaders

img_size = 224 ... batch_size = 32

manual_transforms = transforms.Compose([transforms.Resize((img_size, img_size)), transforms.ToTensor()])

train_dataloader, test_dataloader, class_names = data_setup.create_dataloaders(train_dir=train_dir,
test_dir=test_dir, transform = manual_transforms, batch_size = batch_size)

4: equation 1

height = 224 ... width = 224 ... color_channels = 3 ... patch_size = 16

number_of_patches = int((height * width) / patch_size **2)

class PatchEmbedding(nn.Module):

def __init__(self, in_channels: int=3, patch_size: int=16, embedding_dim: int = 768):

super().__init__()

self.patcher = nn.Conv2d(in_channels=in_channels, out_channels=embedding_dim, kernel_size = patch_size,
stride = patch_size, padding=0)

self.flatten = nn.Flatten(start_dim = 2, end_dim=3)

def forward(self,x):

image_resolution = x.shape[-1] ... assert image_resolution % patch_size == 0, f"invalid size"

x = self.flatten(self.patcher(x))

return x.permute(0,2,1)

5: equation 2, multi-head attention

```
class MultiheadSelfAttentionBlock(nn.Module):
    def __init__(self, embedding_dim:int=768, num_heads:int=12, attn_dropout:float=0):
        super().__init__()
        self.layer_norm = nn.LayerNorm(normalized_shape=embedding_dim)
        self.multihead_attn = nn.MultiheadAttention(embed_dim=embedding_dim, num_heads=num_heads,
                                                    dropout=attn_dropout, batch_first=True)
    def forward(self, x):
        x = self.layer_norm(x)
        attn_output, _ = self.multihead_attn(query=x, key=x, value=x, need_weights=False)
        return attn_output
class MLPBlock(nn.Module):
    def __init__(self, embedding_dim:int=768, mlp_size:int=3072, dropout:float=0.1):
        super().__init__()
        self.layer_norm = nn.LayerNorm(normalized_shape:embedding_dim)
        self.mlp = nn.Sequential(nn.Linear(in_features=embedding_dim, out_features=mlp_size),
                                 nn.GELU(), nn.Dropout(p=dropout), nn.Linear(in_features=mlp_size, out_features=
                                 embedding_dim), nn.Dropout(p:dropout))
    def forward(self,x):
        x = self.mlp(self.layer_norm(x))  ... return x
```

7: create transformer encoder

```
class TransformerEncoderBlock(nn.Module):
    def __init__(self, embedding_dim:int=768, num_heads:int=12, mlp_size:int=3072,
                 mlp_dropout:float=0.1, attn_dropout:float=0):
        super().__init__()
        self.msa_block = MultiheadSelfAttentionBlock(embedding_dim, num_heads=num_heads, attn_dropout=
                                                     attn_dropout)
        self.mlp_block = MLPBlock(embedding_dim=embedding_dim, mlp_size=mlp_size, dropout=mlp_dropout)
    def forward(self,x):
        x = self.msa_block(x)+x  ...  x = self.mlp_block(x)+x
        return x
```

Can also use `torch.nn.TransformerEncoderLayer()`

Pytorch paper replicating continued...

```
class ViT(nn.Module):
    def __init__(self, img_size:int=224, in_channels:int=3, patch_size:int=16, num_transformer_layers:int=12,
                 embedding_dim:int=768, mlp_size:int=3072, num_heads:int=12, attn_dropout:float=0,
                 mlp_dropout:float=0.1, num_classes:int=1000):
        super().__init__()
        assert img_size % patch_size == 0
        self.num_patches = (img_size * img_size) // patch_size**2
        self.class_embedding = nn.Parameter(data=torch.rand(1,1,embedding_dim), requires_grad=True)
        self.position_embedding = nn.Parameter(data=torch.rand(1, self.num_patches+1, embedding_dim), requires_grad=True)
        self.embedding_dropout = nn.Dropout(p=embedding_dropout)
        self.patch_embedding = PatchEmbedding(in_channels=in_channels, patch_size=patch_size, embedding_dim=embedding_dim)
        self.transformer_encoder = nn.Sequential(*[TransformerEncoderBlock(embedding_dim=embedding_dim, num_heads=num_heads,
                                                   mlp_size=mlp_size, mlp_dropout=mlp_dropout)
                                                   for _ in range(num_transformer_layers)])
        self.classifier = nn.Sequential(nn.LayerNorm(normalized_shape=embedding_dim), nn.Linear(in_features=embedding_dim,
                                        out_features=num_classes))
    def forward(self, x):
        batch_size = x.shape[0]
        class_token = self.class_embedding.expand(batch_size, -1, -1)
        x = self.patch_embedding(x) ... x = torch.cat((class_token, x), dim=1) ... x = self.position_embedding + x
        x = self.embedding_dropout(x) ... x = self.transformer_encoder(x) ... x = self.classifier(x[:, 0])
        return x
vit = ViT(num_classes=len(class_names))
```

9: Setting up and training code for ViT model

```
from going_modular.going_modular import engine
optimizer = torch.optim.Adam(params=vit.parameters(), lr=3e-3, betas=(0.9,0.999), weight_decay=0.3)
loss_fn = torch.nn.CrossEntropyLoss()
set_seeds()
results = engine.train(model=vit, train_dataloader=train_dataloader, test_dataloader=test_dataloader,
                       optimizer=optimizer, loss_fn=loss_fn, epochs=10, device=device)
```